Digital Land: Data Center Announcements and the Capture of Community-Created Value in Georgia

Damian Krotov

2026-07-05

Georgia is the fastest-growing new data-center market in the United States, courted with a sales-tax exemption forgoing roughly $2.5 billion a year and backed by nearly 10 GW of newly certified generation. Using a hand-built, source-verified dataset of dated Georgia data-center announcements (2015–2026) and Zillow's ZIP-level home value index, I estimate how announcements capitalize into nearby residential property values with the Callaway–Sant'Anna staggered difference-in-differences estimator. Announcements produce no detectable average effect on ZIP-level home values; pre-trends are flat and the result survives a robustness battery. Read through Henry George's theory of rent, the null is informative: the community-created value of the data center boom is not accruing broadly to nearby homeowners, while its documented fiscal costs — the exemption and ratepayer-backstopped generation for load that may never materialize — are socialized. Land value capture, not equipment subsidies, is the remedy George's framework recommends.

1 Introduction

When Georgia's state auditors examined four large metro-Atlanta data-center projects, the "representative" complex they constructed carried servers and electrical equipment worth about $1.81 billion — sitting on land worth about $26.1 million (S4). That ratio is a modern illustration of the distinction Henry George built *Progress and Poverty* around: capital, which human effort produces and replaces, and land, whose value is created not by its owner but by the community around it (George 1879). Georgia has chosen to subsidize the first — data-center equipment enjoys a sales-tax exemption projected to forgo roughly $2.5 billion in state and local revenue in FY2026, approaching $3 billion in 2027 (S1) — while leaving the second untaxed beyond ordinary rates.

This paper measures one channel through which the data-center boom could create — or destroy — community wealth: the capitalization of project *announcements* into nearby residential property values. An announcement is George's mechanism compressed into a single dated event: news of collective investment (grid capacity, fiber, roads, tax base) arrives, and whatever value that news creates accrues to whoever holds title nearby. I hand-build a source-verified dataset of 48 dated Georgia data-center announcements (2015–2026) totaling over $100 billion in announced investment, match them to 31 ZIP codes, and estimate group-time average treatment effects on Zillow's ZIP-level home value index with the Callaway and Sant'Anna (2021) estimator.

The answer is a null: the simple average post-announcement effect on log home values is 1.2% (95% CI -4.3% to 6.8%), with flat pre-trends (0 of 23 pre-period coefficients significant at 5%). Georgia's homeowners near announced sites have not, on average, received a speculative windfall — nor suffered a broad disamenity discount — at the ZIP level. Section 7 discusses why the null is the analytically useful case for the distributional question: the boom's benefits are not diffusing into neighboring home equity, while its costs (the exemption, and the risk that certified load never materializes) are borne publicly.

2 Background: the boom, the exemption, and the risk

Three institutional facts frame the analysis; each is documented in the repository's source index with the state's own documents.

Scale. As of December 2025 Georgia had an estimated 63 active data centers, 35 under construction, and 249 more announced (S3). On December 19, 2025 the Public Service Commission approved a stipulation certifying 9,885 MW of new generation — approximately 80 percent of it expected to power data centers — while freezing Georgia Power base rates through 2028 and building a menu of remediation measures should contracted load fail to appear (S5).

Subsidy. The state's tax expenditure report puts the combined cost of the data-center and high-technology equipment exemptions at roughly $2.5 billion in FY2026, rising toward $3 billion in FY2027 (S1). The state audit's central counterfactual finding: about 70 percent of Georgia data-center activity would have occurred without the exemption (S2). Every bill that would have curtailed the exemption or strengthened ratepayer protection failed in the 2026 session; the exemption stands until 2032 (S8).

Risk. Nationally, "phantom" interconnection load is documented: Exelon reports only about 22 percent of its 65-GW pipeline through 2040 is likely to materialize (S6), and telemetry analysis by E3 finds fewer than half of studied data centers run above 80 percent of reported peak load (S7). Under the PSC stipulation, Georgia households function as residual insurers of certified load, mitigated by remediation clauses and a Georgia Power backstop through 2031 (S5).

3 The Georgist frame

George's savannah parable places a settler on land with "absolutely no choice" among identical acres; value arrives only as others settle nearby: "He has what, were he in a populous district, would make him rich; but he is very poor." (S10, *Progress and Poverty*, Bk. IV, Ch. 2). Rent, in this account, is community-created, and its private capture is the unearned increment. Four mappings organize the paper: (i) an announcement is community-created value (or damage) arriving as news, testable as capitalization; (ii) the audit's land/equipment split (S4) is George's land–capital distinction in state documents; (iii) the exemption is inverse-George — the public subsidizes capital while leaving site-value gains private; (iv) the stipulation socializes speculative risk, George's critique of land speculation translated to interconnection queues. The remedy George's framework implies — land value taxation / land value capture — taxes the measured uplift (where it exists) without taxing the servers (Plassmann and Tideman 2000).

4 Data

Events (D1). The unit of treatment is a dated public announcement of a data-center project. Each of the 48 verified events carries company, site location, the earliest credible public report date, announced investment, capacity and acreage where stated, current status, and at least one fetched source URL; every row was verified against its source, with judgment calls logged (`data/raw/verification_log.md`). Discovery drew on state and county press, trade press (Data Center Frontier, DCD), local outlets, and development-authority filings. Withdrawn and denied projects are retained — an announcement is news regardless of what later happens, and the "value that never materializes" is part of the research question.

Outcomes (D2). Zillow's ZHVI, ZIP-level, monthly, smoothed and seasonally adjusted; Redfin's ZIP median sale price is a robustness outcome. The panel covers 550 Georgia ZIPs from 2012.

Geography and covariates (D3). Census ZCTA polygons assign each site to a ZIP (stated ZIP when sources give one, point-in-polygon on geocoded coordinates otherwise) and identify first-ring neighbors; ACS 2023 5-year population and median household income enter as covariates. ZIP ZCTA is assumed with the standard caveat.

5 Empirical strategy

Let g index announcement month. I estimate group-time effects $ATT(g,t)$ on log(ZHVI) with the Callaway and Sant'Anna (2021) estimator: not-yet-treated ZIPs as controls (never-treated as robustness — treated ZIPs are compared primarily to future siting corridors, not to systematically different rural ZIPs), three months of permitted anticipation (land-assembly rumors precede press releases), universal base period, covariate adjustment by log population

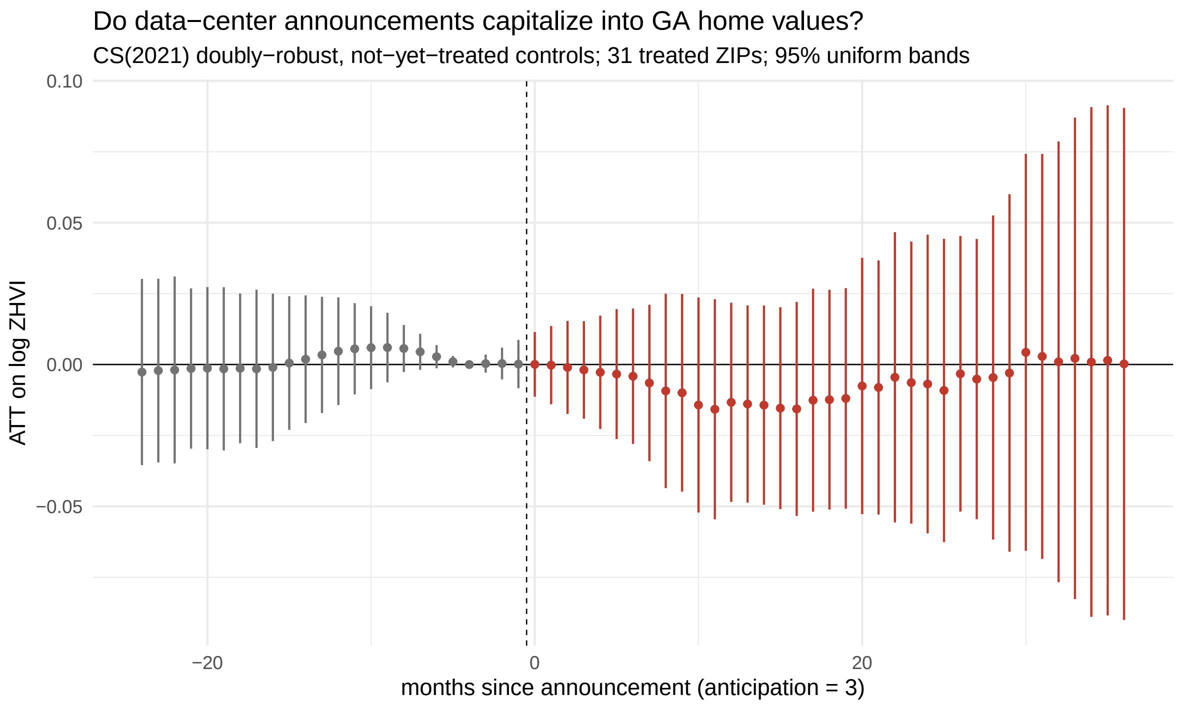

Figure 1: Event-study path of announcement effects on log ZHVI. CS(2021), outcome-regression, not-yet-treated controls, anticipation = 3, 95% bands.

and log income, and clustering by ZIP. Static two-way fixed effects would be contaminated by forbidden comparisons under staggered timing and heterogeneous effects (Goodman-Bacon 2021); the interaction-weighted estimator of Sun and Abraham (2021) is reported as robustness. First-ring neighbor ZIPs are excluded from the control pool (SUTVA hygiene). Design choices follow the practitioner synthesis in Roth et al. (2023).

One pre-registered choice failed mechanically and is documented rather than hidden: the doubly-robust estimator cannot fit its propensity-score component when an announcement cohort contains a single treated ZIP, which is the typical case here; nearly all $ATT(g, t)$ return missing. The outcome-regression estimator retains the covariate adjustment and is the primary specification; doubly-robust without covariates and IPW variants bracket it in the robustness table.

Hypotheses were committed to git before estimation: H1a (premium), H1b (discount), H0 (no detectable effect, consistent with the Virginia evidence of Priest (2026)).

6 Results

The headline estimate is 1.2% on log home value (SE 2.8%; 95% CI -4.3% to 6.8%), across 31 treated ZIPs against 519 controls (109 first-ring neighbors excluded). The dynamic path is flat: pre-period coefficients average 0.12% with 0 of 23 significant at 5%, and post-period estimates oscillate around zero through 36 months. The design can rule out average effects larger than roughly ±5–7 log points; it cannot — and does not claim to — detect hyperlocal effects within a mile of a fence line, which ZIP-level smoothed indices average away.

7 Robustness

spec	ATT (%)	SE (%)	95% CI	n_treated
(0) main: not-yet-treated, anticipation 3	1.24	2.78	[-4.2, 6.7]	31
(a) never-treated controls	1.34	2.85	[-4.3, 6.9]	31
(a') dr, no covariates	3.41	3.59	[-3.6, 10.5]	31
(a'') ipw + covariates	2.71	2.95	[-3.1, 8.5]	31
(b) Sun–Abraham IW, all-horizon post ATT (different estimand — see text)	-0.97	0.28	[-1.5, -0.4]	31
(c) Redfin median sale price outcome	-2.14	3.22	[-8.5, 4.2]	30
(d) drop 2020-01..2021-12	0.35	2.86	[-5.2, 6.0]	31
(e) anticipation = 0	1.01	2.86	[-4.6, 6.6]	31
(e') anticipation = 6	0.98	3.09	[-5.1, 7.0]	31
(g) drop top-3 investment projects	1.34	2.89	[-4.3, 7.0]	28
(h) propensity-trimmed controls	1.1	2.98	[-4.7, 6.9]	31

A permutation placebo reassigning announcement dates uniformly over the observed range (B = 199) yields p = 0.889 for the actual estimate.

Row (b) is not directly comparable to the others: Sun–Abraham's interaction-weighted ATT equal-weights every post horizon (up to +131 months for the earliest cohort) rather than the CS simple aggregation, which is why its estimate is small, negative, and tightly estimated; window-capping it to +36 months produces degenerate singleton cells. Its substantive message matches the rest of the table: no large positive capitalization.

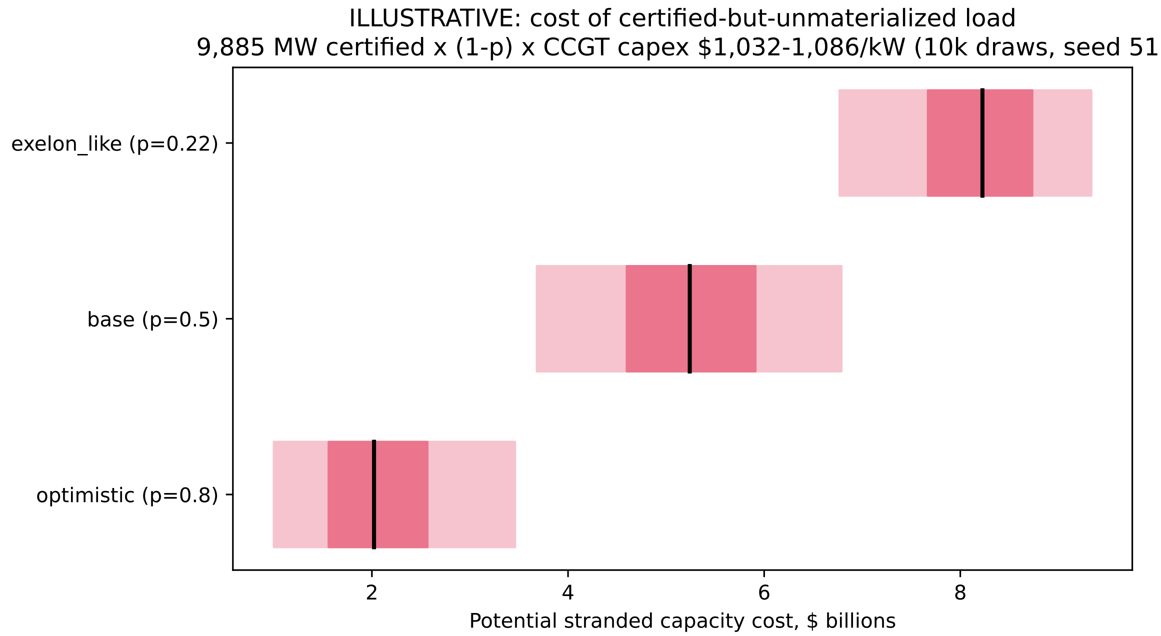

Figure 2: ILLUSTRATIVE stranded-cost simulation (10,000 draws, seed 51). Not a causal estimate.

8 Who bears the risk that the value never arrives?

The capitalization estimate answers who *gains*; the stipulation decides who *insures*. To communicate the asymmetry — not to estimate a causal quantity — an illustrative Monte Carlo prices the certified-but-unmaterialized load: 9,885 MW certified (S5), materialization probability drawn around three scenarios (optimistic 0.8; base 0.5; the Exelon-like 0.22 of S6), and combined-cycle replacement capital cost drawn from the EIA reference range (S9).

Scenario medians and 90% ranges: optimistic: median $2.0B [$1.0, $3.5]; base: median $5.2B [$3.7, $6.8]; exelon_like: median $8.2B [$6.8, $9.3].

Put together: measured private gains to nearby homeowners are approximately zero; the public's exposure — $2.5 billion a year in forgone revenue (S1) plus billions in contingent stranded-capacity risk — is not. George's remedy would invert the arrangement: exempt the $1.81 billion of servers, tax the site value the community creates, and let the public capture the uplift where and when it appears.

9 Limitations

(1) Siting is endogenous; developers choose growth corridors. Not-yet-treated controls, covariate adjustment, and flat pre-trends mitigate but do not prove; conclusions are stated as "capitalization consistent with," never causality established. (2) Announcement dates carry measurement error (rumors precede pressers); the three-month anticipation window absorbs some, and attenuation would bias toward the null. (3) ZIP-level smoothed indices mask hyperlocal effects; a null here does not contradict fence-line testimony, and parcel-level analysis is future work. (4) Adjacent-ZIP spillovers are handled by exclusion, not modeling. (5) Findings are Georgia, 2015–2026, residential; nothing here is investment, legal, or tax advice. (6) AI assistance was used for data collection, source verification, and code under the author's direction and is disclosed per contest rules; every row of the event dataset carries its source, and every statistic in this paper is computed by the released pipeline from `output/results.json`.

References

Callaway, Brantly, and Pedro H. C. Sant'Anna. 2021. "Difference-in-Differences with Multiple Time Periods." *Journal of Econometrics* 225 (2): 200–230.

George, Henry. 1879. *Progress and Poverty: An Inquiry into the Cause of Industrial Depressions and of Increase of Want with Increase of Wealth.* D. Appleton; Company.

Goodman-Bacon, Andrew. 2021. "Difference-in-Differences with Variation in Treatment Timing." *Journal of Econometrics* 225 (2): 254–77.

Plassmann, Florenz, and T. Nicolaus Tideman. 2000. "A Markov Chain Monte Carlo Analysis of the Effect of Two-Rate Property Taxes on Construction." *Journal of Urban Economics* 47 (2): 216–47.

Priest, TODO_VERIFY_first_name. 2026. *TODO_VERIFY Exact Title: Data Centers and Residential Property Values, Virginia Air-Permit DiD.* SSRN Working Paper 6314620.

Roth, Jonathan, Pedro H. C. Sant'Anna, Alyssa Bilinski, and John Poe. 2023. "What's Trending in Difference-in-Differences? A Synthesis of the Recent Econometrics Literature." *Journal of Econometrics* 235 (2): 2218–44.

Sun, Liyang, and Sarah Abraham. 2021. "Estimating Dynamic Treatment Effects in Event Studies with Heterogeneous Treatment Effects." *Journal of Econometrics* 225 (2): 175–99.